

3D Graphics Engine

Team: Mykhailo-Taras Sobko, Anna-Alina Bondarets, Oleksandra Stasiuk

The source code of the project can be found on the GitHub:

https://github.com/MykhailoSobko/3d_graphics_rendering

1. Introduction and aim

This project is about rendering 3D objects. Given one object representation in 3D, the engine gives the opportunity to view it projected onto a 2D screen with the transformations given from the user, such as rotations, zooming, translation, and changing light direction. This will be useful not only for the study purposes, but also as a fundamental step to developing the real game engine.

The project is focused around the real-time computer graphics, or real-time rendering, conceptions. It provides real-time object analysis with an interactive interface through keyboard and mouse. Using the different techniques for rendering, such as ray-tracing and rasterization, the user can respond to rendered images in real time, producing an interactive experience.

2. Rendering problem

Rendering is the final process of creating the actual 2D image or animation from the prepared scene. This can be compared to taking a photo or filming the scene after the setup is finished in real life. Several different, and often specialized, rendering methods have been developed. These range from the distinctly non-realistic wireframe rendering through polygon-based rendering, to more advanced techniques such as: scanline rendering, ray tracing, or radiosity. Rendering may take from fractions of a second to days for a single image/frame. In general, different methods are better suited for either photorealistic rendering, or real-time rendering.

Rendering for interactive media, such as games and simulations, is calculated and displayed in real time, at rates of approximately 20 to 120 frames per second. In real-time rendering, the goal is to show as much information as possible as the eye can process in a fraction of a second (a.k.a. "in one frame": In the case of a 30 frame-per-second animation, a frame encompasses one 30th of a second).

The primary goal is to achieve an as high as possible degree of photorealism at an acceptable minimum rendering speed (usually 24 frames per second, as that is the minimum the human eye needs to see to successfully create the illusion of movement). In fact, exploitations can be applied in the way the eye 'perceives' the world, and as a result, the final image presented is not necessarily that of the real world, but one close enough for the human eye to tolerate.

The rapid increase in computer processing power has allowed a progressively higher degree of realism even for real-time rendering. Real-time rendering is often polygonal and aided by the computer's GPU.

Polygonal modeling is an approach for modeling objects by representing or approximating their surfaces using polygon meshes. The basic object used in mesh modeling is a vertex, a point in three-dimensional space. Two vertices connected by a straight line become an edge. Three vertices, connected to each other by three edges, define a triangle, which is the simplest polygon in Euclidean space. A group of polygons, connected to each other by shared vertices, is generally referred to as an element. Each of the polygons making up an element is called a face.

The flat nature of triangles makes it simple to determine their surface normal. Surface normals are useful for determining light transport in ray tracing, and are a key component of the popular shading models (and the one we use). Note that every triangle has two face normals, which point in opposite directions from each other. In many systems only one of these normals is considered valid – the other side of the polygon is referred to as a backface, and can be made visible or invisible depending on the programmer's desires.

3. Technologies


We decided to develop a real-time rendering model, because of the linear algebra needed in the real-time calculations behind interactions with the object. We used the most popular triangle polygonal approach, which is well-documented and acknowledged in the industry (although voxel rendering, which utilizes regular grids, i. e., cubes, instead of triangles, is becoming more and more popular nowadays) and a flat shading technique, which is relatively simple in technical way compared to others, but utilizes linear algebra methods.

Professional graphic engines may utilize more complex polygons which consist not only of triangles and smarter shading methods, but the aim of this project is to apply linear algebra that is used in the core of the concept of 3D rendering, not some fancy pre-built technologies.

In our engine, meshes are static (they can not change their form in real time using animation), but can be moved (translated, rotated and scaled) relative to the camera. Also, we used a technique with monochrome meshes, where the mesh's single color is chosen at the start, but each polygon is shaded depending on the direction of the light (which, in our realization, can be rotated horizontally in real time). Adding colors to separate polygons or texturing would complicate the project technically (at least, at this stage), but would not bring in more mathematical value.

3.1. Data

Our implementation uses `.obj` files to get the information needed for rendering. A `.obj` file may contain vertex data, free-form curve/surface attributes, elements, free-form curve/surface body statements, connectivity between free-form surfaces, grouping, and display/render attribute information. The most common elements are geometric vertices, texture coordinates, vertex normals, and polygonal faces. Only the information about vertices and faces is needed for our project, as we did not aim to implement complex technical features.



A vertex is specified via a line starting with the letter **v**. That is followed by (x, y, z) coordinates. A right-hand coordinate system is used to specify the coordinate locations. It is used to determine which side of the triangle is the one we need (we do not want to render the inner one). In other words, there are two possible directions of normal vectors for each face, but we need the one directed outside the mesh, not inside. That is why the vertices are specified in a counterclockwise manner, so we can get the direction of the normal from the right-hand rule, which is used in the cross product.

Faces are defined using lists of vertex indices. A face is specified via a line starting with the letter **f**. A valid vertex index matches the corresponding vertex elements of a previously defined vertex list. If an index is positive, it refers to the offset in that vertex list. The compulsory requirement on faces is to be triangular, thus each face should contain three vertices.

An example of a simple pyramid shape represented in such format:

```
v 0 0 0
v 1 0 0
v 1 1 0
v 0 1 0
v 0.5 0.5 1.6
```

```
f 2 1 4
f 2 4 3
f 1 2 5
f 1 5 4
f 4 5 3
f 2 3 5
```

There are 3 possible ways to get such files:

1. It can be written manually for non-complex shapes (for example, we used this for a pyramid (https://github.com/MykhailoSobko/3d_graphics_rendering/blob/main/data/piramid.obj));
2. There are a lot of **.obj** files freely available on the internet. Most of them store additional information for advanced engines. Still, they are valid for our implementation as long as they have the bare minimum information about the vertices and faces described above. With this approach we found the appropriate cube model (https://github.com/MykhailoSobko/3d_graphics_rendering/blob/main/data/cube.obj);

3. Most 3D modeling software like Blender supports exporting objects in `.obj` file format. We used this approach to generate an octahedron model (https://github.com/MykhailoSobko/3d_graphics_rendering/blob/main/data/octahedron.obj) and test it.

Anyone can test our engine using a valid `.obj` file. However, low framerates are possible with complex objects like spheres. We did not aim to make a technically advanced implementation but rather make it from scratch using linear algebra notions.

The `data` folder (https://github.com/MykhailoSobko/3d_graphics_rendering/tree/main/data) in the project root directory includes the described above `.obj` files and each can be freely used.

3.2. OpenGL

OpenGL is a cross-language, cross-platform API for rendering 2D and 3D objects. We used its python implementation for a basic rendering of the triangles. All the transformations were computed by us in the code, as well as the projection onto 2D screen, yet the OpenGL tools were used to draw lines between the points creating a triangle, and to fill it:

```
def draw_triangle(triangle: np.array, color: tuple[float, float, float]) ->
None:
    glColor3f(color[0], color[1], color[2])
    glBegin(GL_TRIANGLE_STRIP)
    glVertex2f(triangle[0][0], triangle[0][1])
    glVertex2f(triangle[1][0], triangle[1][1])
    glVertex2f(triangle[2][0], triangle[2][1])
    glEnd()
```

It also was used to create the initial window and to track the mouse and keyboard inputs:

```
def start(self) -> None:
    """
    Main runner. Initialized and created window,
    proceeds with infinite loop of iterations to
    display and perform transformations from user input.
    """
    # Initialize window configurations
    glutInit()
    glutInitDisplayMode(GLUT_SINGLE | GLUT_RGB)
    glutInitWindowSize(self.window_size[0], self.window_size[1])
    glutInitWindowPosition(self.window_position[0], self.window_position[1])
    self.window = glutCreateWindow(self.app_name)
    self.__on_user_create()

    # Loop begins here
    # Perform object rendering
    glutDisplayFunc(self.__on_user_update)
```

```

# Track input from keyboard and mouse
glutKeyboardFunc(self.__WASD)
glutSpecialFunc(self.__arrows)
glutMouseFunc(self.__mouse)
glutMouseWheelFunc(self.__mouse_wheel)

# Enter OpenGL event processing loop
glutMainLoop()

```

4. Linear algebra applications: theory and implementation

4.1. Homogeneous coordinates

To render a three-dimensional scene on a two-dimensional display screen, we need to determine where on the screen each vertex of the object should be drawn. 3D graphics systems use homogeneous coordinates, or projective coordinates, to project vertices in space.

If homogeneous coordinates of a point are multiplied by a non-zero scalar then the resulting coordinates represent the same point. Since homogeneous coordinates are also given to points at infinity, the number of coordinates required to allow this extension is one more than the dimension of the projective space being considered. Thus, we should use four homogeneous coordinates (x, y, z, w) to specify a point in the projective plane.

To apply the transformations, we should cast a point (x_0, y_0, z_0) from the Euclid coordinates to the homogeneous one by assigning the fourth coordinate to the value $w_0 = 1$. When we want to draw the point, we should transfer it back to the Euclid coordinates by simply normalizing all the coordinates by the w_0 . Thus, when $w_0 \neq 0$, we achieve a point

$\left(\frac{x_0}{w_0}, \frac{y_0}{w_0}, \frac{z_0}{w_0}\right)$ to be drawn after all the transformations and projections.

4.2. Projection matrix

The projection matrix, the same as all the transformation matrices we will be using, is a 4×4 matrix P , determined as follows:

$$P = \begin{pmatrix} a \cdot \frac{1}{\tan(0.5 \cdot fov)} & 0 & 0 & 0 \\ 0 & \frac{1}{\tan(0.5 \cdot fov)} & 0 & 0 \\ 0 & 0 & \frac{far}{far - near} & 1 \\ 0 & 0 & \frac{-far \cdot near}{far - near} & 0 \end{pmatrix},$$


where:

$a = \frac{height}{width}$ is the aspect ratio of the screen dimensions;

$fov = \frac{1}{\tan(0.5 \cdot fov_{rad})}$ is the view frustum with fov_{rad} being the horizontal field-of-view value given in radians;

far is the value defining where to display the object in the back by the *z* (distance from the camera to the far plane);

near is the value defining the front by the *z* (distance from the camera to the near plane);



4.3. Transformation matrices

All the transformation matrices are represented in the four-dimensional homogeneous coordinate system. Then, the transformation applies with the following steps to each single point of a triangle:

1. Cast 3D vector to homogeneous case
2. Multiply transformation matrix by this vector
3. Cast transformed point to 3D

There are the following transformations implemented in the engine:

1. Translation:
$$\begin{pmatrix} 1 & 0 & 0 & T_x \\ 0 & 1 & 0 & T_y \\ 0 & 0 & 1 & T_z \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

@staticmethod

```
def get_translation_matrix(T_x: float, T_y: float, T_z: float) -> np.array:
    return np.array([[1.0, 0.0, 0.0, T_x],
                     [0.0, 1.0, 0.0, T_y],
                     [0.0, 0.0, 1.0, T_z],
                     [0.0, 0.0, 0.0, 1.0]])
```

2. Rotation by x :

$$\begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & \cos(\phi) & -\sin(\phi) & 0 \\ 0 & \sin(\phi) & \cos(\phi) & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

```
@staticmethod
```

```
def get_X_rotation_matrix(degree: float):
    phi = np.deg2rad(degree)
    return np.array([[1.0, 0.0, 0.0, 0.0],
                     [0.0, np.cos(phi), -np.sin(phi), 0.0],
                     [0.0, np.sin(phi), np.cos(phi), 0.0],
                     [0.0, 0.0, 0.0, 1.0]])
```

3. Rotation by y :

$$\begin{pmatrix} \cos(\phi) & 0 & \sin(\phi) & 0 \\ 0 & 1 & 0 & 0 \\ -\sin(\phi) & 0 & \cos(\phi) & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

```
@staticmethod
```

```
def get_Y_rotation_matrix(degree: float):
    phi = np.deg2rad(degree)
    return np.array([[ np.cos(phi), 0.0, np.sin(phi), 0.0],
                     [ 0.0, 1.0, 0.0, 0.0],
                     [-np.sin(phi), 0.0, np.cos(phi), 0.0],
                     [ 0.0, 0.0, 0.0, 1.0]])
```

4. Rotation by z :

$$\begin{pmatrix} \cos(\phi) & -\sin(\phi) & 0 & 0 \\ \sin(\phi) & \cos(\phi) & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

```
@staticmethod
```

```
def get_Z_rotation_matrix(degree: float):
    phi = np.deg2rad(degree)
    return np.array([[np.cos(phi), -np.sin(phi), 0.0, 0.0],
                     [np.sin(phi), np.cos(phi), 0.0, 0.0],
                     [ 0.0, 0.0, 1.0, 0.0],
                     [ 0.0, 0.0, 0.0, 1.0]])
```

5. Scale:

$$\begin{pmatrix} S_x & 0 & 0 & 0 \\ 0 & S_y & 0 & 0 \\ 0 & 0 & S_z & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

```

@staticmethod
def get_scale_matrix(S_x: float, S_y: float, S_z: float) -> np.array:
    return np.array([[S_x, 0.0, 0.0, 0.0],
                     [0.0, S_y, 0.0, 0.0],
                     [0.0, 0.0, S_z, 0.0],
                     [0.0, 0.0, 0.0, 1.0]])

```

4.4. Depth

To draw a correctly projected image of a 3D object, we should care about the order of the triangles to draw. Also, we should not consider the triangles which describe not visible planes to reduce the rendering process.

To find out whether the triangle is visible, we implemented the following steps:

1. Find normal vector of the triangle
2. Calculate the dot product of the norm and a camera-triangle direction
3. If it is positive, then the triangle is not visible

```

@staticmethod
def get_normal(triangle: np.array):
    edge_1 = triangle[1] - triangle[0]
    edge_2 = triangle[2] - triangle[0]

    normal = np.cross(edge_1, edge_2)

    return normal / np.linalg.norm(normal)

normal = self.get_normal(triangle)

if np.dot(normal, triangle[0] - self.camera) > 0:
    return None

```

Moreover, we can use the value of the dot product dp of triangle normal and a light direction to define the depth of the planes. Because vectors of the dot product are normalized, the value of dp will be in the range between -1 and 1 . Thus, we can use its value to define the color scaler, by multiplying on which we can achieve darker colors of the far planes:

```

@staticmethod
def get_color_scaler(dp: float):
    return (1.5 + dp) / 2.5

```

With that filter, we get the value in the range from 0.2 to 1, and will use this number to multiply each component of the original RGB color. By doing that, the less value becomes, the darker color will be. The lower bound is 0.2 and not less for not making it look completely dark (0 stands for black), while the max value of 1 gives us the original color in the light.

Finally, we should render our triangles starting from the farthest ones to avoid the false overlapping, so we created a container to store each visible triangle, and sort it by the sum of z coordinates in each triangle's point. Considering that our *far* and *near* values by z (see section 4.2) are both positive, we should sort the array in the descending order, as the farthest triangle will have the biggest sum of z coordinates.

```
triangles = sorted(triangles,
                    key=lambda tup: tup[0][0][2] + tup[0][1][2] + tup[0][2][2],
                    reverse=True)
```

4.5. Light direction rotation

Considering the original vector of a light direction $(0, 0, -1)$, we can multiply it from left by any transformation matrix to give a new source of light. Considering the case with the rotation matrix, we implemented the horizontal light rotation around the object.



5. Results

We tested our program using some shapes, some of which we got from the internet and others we made ourselves. We also added the screenshots of them with polygon coloring turned off, so only the edges of triangles are visible. On the last page of this document you can find the screenshots of the rendering of our predefined objects, but for the animated demonstration you can check the *demo* folder in our GitHub repository:

https://github.com/MykhailoSobko/3d_graphics_rendering/tree/main/demo

The results in the final version of our project represented the objects as they should, although we encountered some severe visual bugs during development. At first, we did not scale the object correctly before the projection onto the screen, so the output was incorrect. Then, we made a mistake in calculating whether faces were visible to the camera, which resulted in incorrect face rendering. After fixing that, all simple objects began to render correctly. However, the problem remained for complex meshes, where some faces, which should have been hidden from the camera by the object's other parts, were incorrectly rendered in front of them. This bug was due to a mistake in our approach to sorting triangles before rendering. After we dealt with this problem, the results were satisfying.

We decided to render a spaceship model to compare our results to other rendering software because of its geometry and representability. To the left is our render, and to the right is a render by Windows 3D viewer.



A noticeable difference is the FOV (field of view). Our model appears closer than Microsoft's implementation. We would have to adjust our engine's parameters to get a different perspective.

Even though our engine uses only basic rendering techniques, it is slower in terms of performance than other industry equivalents. Usually, professionals use more technically advanced tools because rendering requires a lot of processing power.

6. Literature

1. <https://canvas.projekti.info/ebooks/Mathematics%20for%203D%20Game%20Programming%20and%20Computer%20Graphics,%20Third%20Edition.pdf>
2. A tutorial-blog
<http://blog.wolfire.com/2009/07/linear-algebra-for-game-developers-part-1/> (as well as the next parts)
3. Tutorials on PyOpenGL:
<https://dev.to/suvink/graphic-designing-using-opengl-and-python-beginners-1i4f>
<https://pythonprogramming.net/opengl-rotating-cube-example-pyopengl-tutorial/>

Official documentation:

<http://pyopengl.sourceforge.net/documentation/manual-3.0/index.html>

